Microsoft Corporation

Q1 2026

Share

10/30/2025

Operator

Conference Operator

Greetings and welcome to the Microsoft Fiscal Year 2026 First Quarter Earnings Conference Call. At this time, all participants are in a listen-only mode. A question and answer session will follow the formal presentation. If anyone should require operator assistance, please press star zero on your telephone keypad. As a reminder, this conference is being recorded. It is now my pleasure to introduce Jonathan Nielsen, Vice President of Investor Relations. Please go ahead.

Jonathan Nielsen

Vice President of Investor Relations

Good afternoon and thank you for joining us today. On the call with me are Satya Nadella, Chairman and Chief Executive Officer, Amy Hood, Chief Financial Officer, Alice Jolla, Chief Accounting Officer, and Keith Dolliver, Corporate Secretary and Deputy General Counsel. On the Microsoft Investor Relations website, we will provide our earnings press release and financial summary slide deck, which is intended to supplement our prepared remarks and provides the reconciliation of differences between GAAP and non-GAAP financial measures. More detailed outlook slides will be available on the Microsoft Investor Relations website. On this call, we will discuss certain non-GAAP items. the non-GAAP financial measures provided should not be considered as a substitute for or superior to the measures of financial performance prepared in accordance with GAAP. They are included as additional clarifying items to aid investors in further understanding the company's first quarter performance in addition to the impact these items and events have on the financial results. All growth comparisons we make on the call today relate to the corresponding period of last year unless otherwise noted. will also provide growth rates in constant currency when available as a framework for assessing how our underlying

business performed excluding the effect of foreign currency rate fluctuations where growth rates are the same in constant currency we will refer to the growth rate only we will post our prepared remarks to our website today's call is being recorded if you ask a question it will be included in our live transmission in the transcript and in any future use of the recording. You can replay the call and view the transcript on the Microsoft Investor Relations website. During this call, we will be making forward-looking statements, which are predictions, projections, or other statements about future events. These statements are based on current expectations and assumptions that are subject to risks and uncertainties. Actual results could materially differ because of factors discussed in today's earnings press release in the comments made during this conference call, and in the risk factor section of our Form 10-K, Forms 10-Q, and other reports and filings with the Securities and Exchange Commission. We do not undertake any duty to update any forward-looking statement. And with that, I'll turn the call over to Satya.

Satya Nadella
Chairman and Chief Executive Officer

Thank you, Jonathan. It was a very strong start to our fiscal year. Microsoft cloud revenue surpassed $49 billion, up 26% year over year. And our commercial RPO increased over 50% to nearly $400 billion with a weighted average duration of only two years. We're seeing increasing demand and diffusion of our AI platform and family of co-pilots, which is fueling our investments across both capital and talent. When it comes to infrastructure, we're building a planet-scale cloud and an AI factory, maximizing tokens per dollar per watt, while supporting the sovereignty needs of customers and countries. We're innovating rapidly across the family of co-pilots, spanning the high value domains of information work, coding, security, science, health, and consumer. And as you saw yesterday, we closed a new definitive agreement with OpenAI, marking the next chapter in what is one of the most successful partnerships and investments our industry has ever seen. This is a great milestone for both companies, and we continue to benefit mutually from each other's growth across multiple dimensions. Already, we have roughly 10x our investment. OpenAI has contracted an incremental $250 billion of Azure services. Our RevShare exclusive IP rights and API exclusivity for Azure continue until AGI or through 2030. And we have extended the model and product IP rights through 2032. And we are also energized to innovate and pursue AI advancements with both talent and compute investments that have real-world impact. With that, let's turn to our momentum across our AI platform and co-pilots, as well as with agents. We have the most expansive data center fleet for the AI era, and we are adding capacity at an unprecedented scale. We will increase our total AI capacity by over 80% this year and roughly double our total data center footprint over the next two years, reflecting the demand signals we see. Just this quarter, we announced the world's most powerful AI data center, Fairwater in Wisconsin, which will go online next year and scale to 2 gigawatts alone. And we have deployed the world's first large-scale cluster of NVIDIA GB300s. We are building a fungible fleet that's being continuously modernized

and spans all stages of the AI lifecycle, from pre-training to post-training to synthetic data generation and inference, And it also goes beyond gen AI workloads to recommendation engines, databases, and streaming. We're optimizing this fleet across silicon systems and software to maximize performance and efficiency. It's this combination of fungibility and continuous optimization that allows us to deliver the best ROI and TCO for us and our customers. For example, during the quarter, we increased the token throughput for GPT-4.1 and GPT-5, two of the most widely used models by over 30% per GPU. We also have the most comprehensive digital sovereignty platform. Azure customers in 33 countries are now developing their own cloud and AI capabilities within their borders to meet local data residency requirements. In Germany, for example, OpenAI and SAP will rely on Azure to deliver new AI solutions to the public sector. On top of this infrastructure, we're building Azure AI Foundry to help customers build their own AI apps and agents. We have 80,000 customers, including 80% of the Fortune 500. We offer developers and enterprise access to over 11,000 models more than any other vendor, including as of this quarter, OpenAI's GPT-5 as well as XAI's Grok 4. For example, Ralph Lauren used Foundry to build conversational shopping experience in its app, enabling customers to describe what they're looking for and get personalized recommendations. And OpenEvidence used Foundry to create its AI-powered clinical assistant, which surfaces relevant medical information to physicians and helps streamline charting. When it comes to our first-party models, we are excited by the performance of our new MAI models for text, voice, and image generation, which debuted among the top in the industry leaderboards. And we continue to make great progress with our five family of SLMs, which now have been downloaded over 60 million times, up 3x year over year. Beyond Models and Foundry, we are providing everything developers need to design, customize, and manage AI applications and agents at scale. Our new Microsoft Agent Framework helps developers orchestrate multi-agent systems with compliance, observability, and deep integration out of the box. For example, KPMG used the framework to modernize the audit process, connecting agents to internal data with enterprise-grade governance and observability. These kinds of real production-scale AI deployments are driving Azure's overall growth. And once again, this quarter, Azure took share. Now let's turn to applications and agents we ourselves are building on this platform. We now have 900 million monthly active users of our AI features across our products. And our first-party family of co-pilots now has surpassed 150 million monthly active users across information work, coding, security, science, health, and consumer. When it comes to information work, we continue to innovate with Microsoft 365 Copilot. Copilot is becoming the UI for the agentic AI experience. We have integrated chat and agentic workflows into everyday tools like Outlook, Word, Excel, PowerPoint, and Teams. Just nine months since release, tens of millions of users across Microsoft 365 customer base are already using chat. Adoption is accelerating rapidly, growing 50% quarter over quarter, and we continue to see usage intensity increase. This quarter, we also introduce agent mode, which turns single prompts into expert quality Word documents, Excel spreadsheets, PowerPoint presentation, and then iterate to deliver the final product, much like agent mode in coding tools today. We're thrilled by the early response, including third-party benchmarks that rank it best in class. Beyond individual productivity,

Copilot is multiplayer. With Teams mode announced this week, you can now invite colleagues into a Copilot conversation. And our collaborative agents like facilitator and project manager prep meeting agendas, take notes, capture decision, kick off group tasks. We are seeing a growing Copilot agent ecosystem with top ISVs like Adobe, Asana, Jira, LexisNexis, SAP, ServiceNow, Snowflake, and Workday all building their own agents that connect to Copilot. And customers are also building agents for their mission-critical business processes and workflows using tools like Copilot Studio and integrating them into Copilot. The overall number of agent users doubled quarter over quarter. And just yesterday, we announced App Builder, a new copilot agent that lets anyone create and deploy task-specific apps and agents in minutes grounded in Microsoft 365 context. All this innovation is driving our momentum. Customers continue to adopt Microsoft 365 Copilot at a faster rate than any other new Microsoft 365 suite. All up, more than 90% of the Fortune 500 now use Microsoft 365 Copilot. Accenture, Bristol-Myers Squibb, EY Global, and the UK's Tax and Payment and Customs Authority all purchased over 15,000 seats this quarter. Lloyds Banking Group has deployed 30,000 seats, saving each employee an average of 46 minutes daily. and a large majority of our enterprise customers continue to come back to purchase more seats. Our partner, PwC alone added 155,000 seats this quarter, and now has over 200,000 deployed across its global operations. In just six months, PwC employees interacted with Microsoft 365 Co-Pilot over 30 million times, and they credit this agentic transformation with saving millions of hours in employee productivity. When it comes to coding, GitHub Copilot is the most popular AI pair programmer now with over 26 million users. For example, tens of thousands of developers at AMD use GitHub Copilot, accepting hundreds of thousands of lines of code suggestions each month and crediting it with saving months of development time. All up, GitHub is now home to over 180 million developers and the platform is growing at the fastest rate in its history, adding a developer every second. 80% of new developers on GitHub start with Copilot within the first week. Overall, the rise of AI coding agents is driving record usage with over 500 million pull requests merged over the past year. And just yesterday at GitHub Universe, we introduced Agent HQ. GitHub Copilot and Agent HQ is the organizing layer for all coding agent extending the GitHub primitives like PRs, issues, actions to coding agents from OpenAI, Anthropic, Google, Cognition, XAI, as well as OSS and in-house models. GitHub now provides a single mission control to launch, manage, and review these agents, each operating from its own branch with built-in controls, observability, and governance. We're building a similar system in security with over three dozen agents in co-pilot integrated across Entra, Defender, Purview, and Intune. For example, with our phishing triage agent in Defender, studies show that analysts can be up to 6.5 times more efficient in detecting malicious mails. In health, Dragon Co-Pilot helps providers automate critical workflows. This quarter alone, we helped document over 17 million patient encounters, up nearly 5x year over year. More than 650 healthcare organizations have purchased our ambient listening tech to date, including University of Michigan Health, where over 1,000 physicians are actively using it. Finally, when it comes to AI consumer experiences, we are excited about all the progress Copilot is making, starting with Windows. Every Windows 11 PC now is an AI PC. Just two weeks ago, we introduced new ways to speak naturally to

your computer, including a co-pilot wake word. With Vision, co-pilot sees what you see on your screen and you can have a real-time conversation about it. And with Action, it takes real action on your behalf, interacting with both web and desktop apps. In Edge, we are introducing first of its kind AI features to automate multi-step workflows within the browser and help you pick up right where you left off. Edge now has taken share for 18 consecutive quarters. In Bing, our overview pages now include embedded conversational capabilities. We took share again in search. And daily users of our Copilot consumer app increased nearly 50% quarter over quarter. Among many updates we made last week is Groups, which turns Copilot for the first time into a shared experience. We also are creating a great consumer subscription offer with Microsoft 365 Premium. It brings together our office applications and advanced co-pilot features with high usage limits, giving individuals the flexibility to bring their own AI to work in a secure way. Finally, in gaming, Copilot provides a voice-first immersive experience across PC, mobile, and our new Xbox ally. Beyond our family of Copilots and AI platform, we are seeing strong momentum across the portfolio. Cloud migrations are accelerating. In data and analytics, fabric revenue grew 60%, which is faster than any other data and analytics platform in the industry. We now have 28,000 paid fabric customers. In databases, SQL DB, Hyperscale revenue was up nearly 75%, 50% in Cosmos DB. In business applications, Dynamics 365 gained share. In security, our end-to-end stack is now informed by 100 trillion daily signals. 1 billion monthly active users of Entra, 16 billion co-pilot interactions audited by Purview, up 72% quarter over quarter. 40,000 Sentinel customers, and we took share across all categories we serve in security. In LinkedIn, nearly 1.3 billion members. And finally, in gaming, we expanded our reach across every endpoint focused on our high margin content and services. We launched critically acclaimed games like Keeper, Ninja Gaiden 4, and Outer Worlds 2, reaching 155 million monthly active users in Minecraft, an all-time high, and set new record for overall content and services revenue for the quarter. We also saw a great response to Xbox Ally launch two weeks ago and set new records for players on PCs. In closing, our planet-scale cloud and AI factory together with co-pilots across high-value domains is driving broad diffusion and real-world impact. And we continue to increase our investments in AI across both capital and talent to meet the massive opportunity ahead. With that, let me turn it over to Amy to walk through our financial results and outlook, and I look forward to rejoining for your questions.

Amy Hood
Chief Financial Officer

Thank you, Satya, and good afternoon, everyone. First, as you heard from Satya, we were pleased to announce the next phase of our partnership with OpenAI yesterday. They continue to choose Microsoft to power their workloads. And together, we remain committed to driving innovation that meets real-world needs. Our Q1 results were not impacted by the deal signed this week. Now, on to the quarter. We delivered a strong start to our fiscal year, exceeding expectations across revenue, operating income, and earnings per share. We also saw continued share gains across many of our businesses, demonstrating

our leadership position in key markets. This quarter, Revenue was $77.7 billion, up 18% and 17% in constant currency. Gross margin dollars increased 18% and 16% in constant currency, while operating income increased 24% and 22% in constant currency. And earnings per share was $4.13, an increase of 23% and 21% in constant currency when adjusted for the impact of our investments in OpenAI. FX impact was roughly in line with guidance. Company gross margin percentage was 69%, down slightly year over year, driven by investments in AI, including the impact of scaling our AI infrastructure and the growing usage of our AI product features. This was partially offset by ongoing efficiency gains, particularly in Azure and M365 commercial cloud. Operating expenses increased 5% and 4% in constant currency, driven by investments in cloud and AI engineering, including compute capacity and AI talent to support product development across the portfolio. Operating margins increased year over year to 49% and we're ahead of expectations with stronger than anticipated results in high margin businesses this quarter. When adjusted for the impact from our investments in OpenAI, other income and expense was $401 million as interest income more than offset interest expense, which includes the interest payments related to data center finance leases. Capital expenditures were $34.9 billion driven by growing demand for our cloud and AI offerings. This quarter, Roughly half of our spend was on short-lived assets, primarily GPUs and CPUs, to support increasing Azure platform demand, growing first-party apps and AI solutions, accelerating R&D by our product teams, as well as continued replacement for end-of-life server and networking equipment. The remaining spend was for long-lived assets that will support monetization for the next 15 years and beyond, including $11.1 billion of finance leases that are primarily for large data center sites. And cash paid for PP&E was $19.4 billion. As a reminder, the difference between total capex and cash paid for PP&E is primarily due to finance leases, as well as the normal timing of goods received but not yet paid. Cash flow from operations was $45.1 billion, up 32%, driven by strong cloud billings and collections, partially offset by higher supplier payments. And free cash flow increased 33%, to $25.7 billion with minimal impact from a sequential increase in CapEx given the higher mix of finance leases. And finally, we returned $10.7 billion to shareholders through dividends and share repurchases. Now to our commercial results. Commercial bookings increased 112% and 111% in constant currency and were significantly ahead of expectations driven by Azure commitments from OpenAI as well as continued growth in the number of $100 million plus contracts for both Azure and M365. These results do not include any impact from the incremental $250 billion Azure commitments from OpenAI announced yesterday. Commercial remaining performance obligation increased to $392 billion and was up 51% year over year. The balance has nearly doubled over the past two years. And even with this growth, our weighted average duration has been relatively stable at approximately two years. Microsoft Cloud revenue was $49.1 billion ahead of expectations and grew 26% and 25% in constant currency. Microsoft Cloud gross margin percentage was slightly better than expected at 68%, and down year over year due to the investments in AI that were partially offset by ongoing efficiency gains as noted earlier. Now to segment results. Revenue from productivity and business processes was $33 billion and grew 17% and 14% in constant currency. M365 commercial cloud revenue increased 17% and 15% in constant currency with one

point of benefit from in-period revenue recognition. Year-over-year growth was driven by both ARPU and seats, with ARPU growth again led by E5 and M365 Copilot. Paid M365 commercial seats grew 6% year-over-year, with installed base expansion across all customer segments, though primarily in our small and medium businesses and frontline worker offerings. M365 commercial products revenue increased 17% and 14% in constant currency ahead of expectations, due to higher than expected Office 2024 transactional purchasing. M365 consumer cloud revenue increased 26% and 25% in constant currency, again driven by ARPU growth. M365 consumer subscriptions grew 7% to over $90 million. LinkedIn revenue increased 10% and 9% in constant currency, driven by marketing solutions. The talent solutions business was impacted by continued weakness in the hiring market. Dynamics 365 revenue increased 18% and 16% in constant currency with continued growth across all workloads. Segment gross margin dollars increased 19% and 16% in constant currency, and gross margin percentage increased, driven by efficiency gains in M365 commercial cloud that were partially offset by investments in AI, including the impact of growing usage in M365 co-pilot chat. Operating expenses increased 6% and 5% in constant currency, and operating income increased 24% and 20% in constant currency. Operating margins increased three points year over year to 62%, driven by the higher gross margin noted earlier, as well as improved operating leverage. Next, the intelligent cloud segment. Revenue was $30.9 billion and grew 28% and 27% in constant currency. In Azure and other cloud services, where we continue to see accelerating demand, Revenue grew 40% and 39% in constant currency. Results were ahead of expectations driven by better than expected growth in

Q1 MSFT 2026 Q&A
0:00 58:40
15 30 1.0x

that were partially offset by efficiency gains in Azure. Operating expenses increased 4% and operating income grew 27%. Operating margins were 43%, down only slightly year over year, as increased investments in AI were mostly offset by improved operating leverage. Now to more personal computing. Revenue was $13.8 billion and grew 4%. Windows OEM and devices revenue increased 6% year over year, significantly ahead of expectations, driven by strong demand ahead of Windows 10 end of support, as well as a benefit from inventory levels that remain elevated. Search and news advertising revenue XTAC increased 16% and 15% in constant currency, driven by growth in volume as well as a continued benefit from third-party partnerships that was better than expected. And in gaming, revenue decreased 2% and 3% in constant currency. Against a strong prior year comparable, Xbox content and services revenue increased 1% and was relatively unchanged in constant currency, driven by better-than-expected performance from third-party content. Segment gross margin dollars increased 11% and 10% in constant currency, and gross margin percentage increased year over year, driven by sales mix shift to higher margin businesses. Operating expenses increased 4% and 3% in constant currency,

and operating income increased 18% and 16% in constant currency. Operating margins increased three points year over year to 30%, driven by the higher gross margin noted earlier. Now, moving to our Q2 outlook, which unless specifically noted otherwise, is on a US dollar basis. Based on current rates, we expect FX to increase total revenue growth by two points. Within the segments, we expect FX to increase revenue growth by two points in productivity and business processes and intelligent cloud, and one point in more personal computing. We expect FX to increase COGS and operating expense growth by one point, starting with the total company. We expect revenue of 79.5 to 80.6 billion U.S. dollars or growth of 14 to 16 percent. We expect COGS of 26.35 to 26.55 billion U.S. dollars or growth of 21 to 22 percent and operating expense of 17.3 to 17.4 billion U.S. dollars or growth of 7 to 8 percent. Operating margins should be relatively flat year over year and down sequentially aligned with historic seasonality. Now, other income and expense. The combination of OpenAI's conversion to a public benefit corp and the ongoing nature of our partnership will result in increased volatility. Therefore, going forward, we'll provide our outlook excluding any impact from our investments in OpenAI. On that basis, in Q2, other income and expense is estimated to be roughly $100 million as interest income will more than offset interest expense. And we expect our Q2 effective tax rate to be approximately 19%. next capital expenditures with accelerating demand and a growing rpo balance we're increasing our spend on gpus and cpus therefore total spend will increase sequentially and we now expect the fy26 growth rate to be higher than fy25 as a reminder there can be quarterly spend variability from cloud infrastructure buildouts and the timing of delivery of finance leases Next, our commercial business. In commercial bookings, we expect healthy growth in the core business on a low expiry base when adjusted for the OpenAI contracts in the prior year. And we expect commercial bookings will be positively impacted by the significant OpenAI commitments announced yesterday. As a reminder, larger long-term Azure contracts, which are more unpredictable in their timing, drive increased quarterly volatility in our bookings growth rate. Microsoft Cloud grows margin percentage to be roughly 66% down year over year, driven by the continued investments in AI, as well as the McShift to Azure. Now to segment guidance. In productivity and business processes, we expect revenue of 33.3 to 33.6 billion US dollars or growth of 13 to 14%. In M365 Commercial Cloud, we expect revenue growth to be between 13 and 14% in constant currency. with business trends that remain relatively stable quarter over quarter. ARPU growth will again be driven by E5 and M365 co-pilot. M365 commercial products revenue growth should be in the low to mid-single digits. As a reminder, M365 commercial products includes components that can be variable due to the end-period revenue recognition dynamics. M365 consumer cloud revenue growth should be in the mid-20s driven by growth in ARPU. For LinkedIn, we expect revenue growth of approximately 10%. And in Dynamics 365, we expect revenue growth to be in the mid to high teens with continued growth across all workloads. For Intelligent Cloud, we expect revenue of 32.25 to 32.55 billion US dollars or growth of 26 to 27%. In Azure, we expect Q2 revenue growth of approximately 37% in constant currency as demand remains significantly ahead of the capacity we have available. And while we're accelerating the amount of capacity we're bringing online, we will continue to balance Azure revenue growth with the

growing needs across our first-party apps and AI solutions, our own R&D efforts, and the end-of-life server replacements. Therefore, we now expect to be capacity constrained through at least the end of our fiscal year. As a reminder, there can be quarterly variability in the year-on-year growth rates depending on the timing of capacity delivery and when it comes online. as well as from in-prem revenue recognition, depending on the mix of contracts. In our on-premises server business, we expect revenue to decline in the low to mid single digits with ongoing customer shift to cloud offerings. In more personal computing, we expect revenue to be in the 13.95 to 14.45 billion US dollars. Windows OEM and devices revenue should decline in the mid single digits. We expect continued momentum from Windows 10 end of support, although growth rates will be impacted by elevated inventory levels at the end of Q1 that we expect to come down through the quarter. Therefore, Windows OEM revenue should decline low to mid-single digits. The range of potential outcomes remains wider than normal. Devices revenue should decline year over year. Search and use advertising XTAC revenue growth should be in the low double digits, down sequentially as growth rates normalize, following the benefit from third-party partnerships noted earlier. Growth will continue to be driven by volume and revenue per search across Edge and Bing. And in Xbox content and services, we expect revenue to decline the low to mid single digits against a prior year comparable that benefited from strong first-party performance, partially offset by growth in subscriptions. And hardware revenue should decline year over year. And in closing, demand signals across bookings, RPO, and product usage are accelerating faster than we expected. We're investing in infrastructure, AI talent, and product innovation to capture that momentum and expand our leadership position. And we remain focused on delivering real value to our customers that results in durable revenue growth for the long term. With that, let's go to Q&A, Jonathan.

Jonathan Nielsen
Vice President of Investor Relations

Thanks, Amy. We'll now move over to Q&A. Out of respect for others on the call, we request that participants please only ask one question. Operator, can you please repeat your instructions?

Operator
Conference Operator

Ladies and gentlemen, if you would like to ask a question, please press star 1 on your telephone keypad and the confirmation tone will indicate your line is in the question queue. You may press star 2 if you would like to remove your question from the queue. For participants using speaker equipment, it may be necessary to pick up your handset before pressing the star keys. And our first question comes from the line of Keith Weiss with Morgan Stanley. Please proceed.

Keith Weiss
Analyst, Morgan Stanley

Excellent. Thank you guys for taking the question. And congratulations on another outstanding quarter. And if I'm looking at Microsoft, this is two quarters in a row. We're really seeing results that are well ahead of anybody's expectations when we're thinking about this company a year ago or five years ago. 111% commercial bookings growth was not on anybody's bingo card, if you will. Yet the stock is underperforming the broader market. And the question I have is kind of getting at the zeitgeist that I think is weighing on the stock. And it's something about to change. And I think AGI is kind of a nomenclature, a shorthand for that. And it's something that's still included in your guys' OpenAI agreements. So, Sakia, when we think about AGI or we think about how application and computing architectures are changing, is there anything that you see on the horizon, whether it's AGI or something else, that could potentially change what appears to be a really strong positioning for Microsoft in the marketplace today where that strength will perhaps weaken on a go-forward basis? Is there anything that you're worrying about in that evolution, and particularly the evolution of these generative AI models?

Satya Nadella
Chairman and Chief Executive Officer

Thank you, Keith, for the question. So here's how I'd say that. I think there are two parts. We feel very, very good about even this, I'd say, the new agreement that we now have with OpenAI because I think even it just creates more certainty to all of the IP relationship we have as it relates to even this definition of AGI. But beyond that, I think your question touches on something that's pretty important, which is how are these AI systems going to truly be deployed in the real world and make a real difference and make a return for both the customers who are deploying them and then obviously the providers of these systems? And I think the best way to characterize the situation is that even as the intelligence capability increases, let's even say exponentially, like model version over model version, The problem is it's always going to still be jagged, right? I think the term people use is the jagged intelligence, or spiky intelligence, right? So you may even have a capability that's fantastic at a particular task, but it may not uniformly grow. So what is required is, in fact, these systems, whether it is GitHub Agent HQ or the M365 Copilot system. Don't think of this as a product. Think of it as a system that in some sense smooths out those jagged edges and really helps the capability. I mean, just to give you a flavor for it, right? So if I am in M365 Copilot, I can generate an Excel spreadsheet. The good news is now an expert spreadsheet does understand OfficeJS, has the formulas in it. It feels like, wow, it is a great spreadsheet created by a good modeler. The more interesting thing is I can go into agent mode in Excel and iterate on that model, and yet it will stay on Rails. It won't go off Rails. It will be able to do the iteration. Then I can even give it to the analyst agent, and then it will even make sense of it like a data analyst would of an Excel model. The reason I say all of that is because that's the type of construction that will be needed, even when the model is magical, all-powerful. I think we will be in this jagged intelligence phase for a long time. So

one of the fundamental things that these, whether it's GitHub, whether it's security, whether it's M365, the three main domains we're in, we feel very, very good about building these as organizing layers for agents to help customers. And by the way, that's the same thing that we want to put into Foundry for our third-party customers. So that's kind of how people will build these multi-agent systems. So I feel actually pretty good about both the progress in AI. I don't think AGI, as defined at least by us in our contract, is ever going to be achieved anytime soon. but I do believe we can drive a lot of value for customers with advances in AI models by building these systems. So it's kind of the real question that needs to be well understood, and I feel very, very confident about our ability to make progress.

Keith Weiss
Analyst, Morgan Stanley

Excellent. That's super helpful.

Jonathan Nielsen
Vice President of Investor Relations

Thanks, Keith. Operator, next question, please.

Operator
Conference Operator

The next question comes from the line of Brent Thill with Jefferies. Please proceed.

Brent Thill
Analyst, Jefferies

Thanks. Amy, on the bookings blowout, I guess many are, you know, somewhat concerned about concentration risk. And I think you noted a number of $100 million contracts, not to go into a lot of detail, but can you just give us a sense of what you're seeing in that 51% RPO and 110 plus percent bookings growth that gives you confidence about what you're seeing in terms of the breadth and, you know, extent of some of these deals on a global basis. Thanks.

Amy Hood
Chief Financial Officer

Thanks, Brent. A couple things to maybe take a step back on RPO. With a nearly $400 billion balance, we've been trying to help people understand sort of how to think about really the breadth of that. It covers numerous products. It covers

customers of all sizes. And that's been a balance that we've been growing, obviously, at a good clip. But what people need to realize is it sits across multiple products because of the things Satya's talking about around creating systems and where we're investing. And if you're going to have that type of balance, and then more importantly, if have the weighted average duration be two years, it means that most of that is being consumed in relatively short order. People are not consuming, and I say this broadly, unless there's value. And I think this is why we keep coming back to are we creating real-world value in our AI platforms, in our AI solutions and apps and systems. And so I think sort of the way to think about RPO is it's been building across a number of customers, We're thrilled to have OpenAI be a piece of that. We're learning a ton and building leading systems because of it that are being used at scale that benefits every other customer. And so it's why we've tried to give a little bit more color to that RPO balance because I do understand that there have been a lot of concerns or questions about is it long dated? Is it coming over a long period of time? And hopefully this is helpful for people to realize that these are contracts being signed by customers who intend to use it in relatively short order. And at that type of scale, I think that's a pretty remarkable execution.

Jonathan Nielsen
Vice President of Investor Relations

Thank you. Thanks, Brent. Operator, next question, please.

Operator
Conference Operator

The next question comes from the line of Mark Mortler with Bernstein Research. Please proceed.

Mark Mortler
Analyst, Bernstein Research

Thank you very much for taking my question, and congratulations on the quarter. It's pretty amazing what you guys are doing. Satya and Amy, I'd like to ask you the number one question I receive, whether from investors or at AI conferences I attend. How much confidence do you have that the software, even the consumer internet business, can monetize all the investments we're seeing globally, or frankly, are we in a bubble? In fact, Amy, what would be the factors you'd be watching to it? to assure that you're not overbuilding the current demand and that demand will sustain. Thank you.

Amy Hood
Chief Financial Officer

Maybe I'll start, Satya, and then you could add. Let me talk a little bit about maybe connecting a couple of the dots, because with $400 billion of RPO that's sort of short-dated, as we talked about, our needs to continue to build out the infrastructure is very high. And that's for booked business today. That is not any new booked business we started trying to accomplish on October 1st, right? And so the way to think about that, and you saw it this quarter in particular, and as we talked about 26, the remainder, number one, we're pivoting toward increasingly, we talked about this, short-lived assets, both GPUs and CPUs. Again, we talk about all these workloads are burning both in terms of app building. Now, when that happens, short-lived assets generally are done to match sort of the duration of the contracts or the duration of your expectation of those contracts. And so I sometimes think when people think about risk, they're not realizing that most of the lifetimes of these and the lifetimes of the contracts are very similar. And so when you think about having revenue and the bookings and coming on the balance sheet and the depreciation of short-lived assets, they're actually quite matched, Mark. And as you know, we've spent the past few years not actually being short GPUs and CPUs per se. We were short the space or the power is the language we use to put them in. So we spent a lot of time building out that infrastructure. Now we're continuing to do that, also using leases. Those are very long-lived assets, as we've talked about, 15 to 20 years. And over that period of time, do I have confidence that we'll need to use all of that? It is very high. And so when I think about sort of balancing those things, seeing the pivot to GPU-CPU short-lived, seeing the pivot in terms of how those are being utilized, we are, and I said this now, we've been short-lived. Now, for many quarters, I thought we were going to catch up. We are not. Demand is increasing. It is not increasing in just one place. It is increasing across many places. We're seeing usage increases in products. We are seeing new products launch that are getting increasing usage and increasing usage very quickly. You know, when people see real value, they actually commit real usage. And I sometimes think this is where this cycle needs to be thought through completely is that, you know, when you see these kind of demand signals and we know we're behind, we do need to spend. But we're spending with a different amount of confidence in usage patterns and in bookings, and I feel very good about that. I have said we are now likely to be short capacity to serve the most important things we need to do, which is Azure, our first-party applications. We need to invest in product R&D, and we're doing end-of-life replacements in the fleet. So we're going to spend to make sure that happens. It's about modernization. It's about high quality. It's about service delivery, and it's about meeting demand. And so I feel good about doing that, and I feel good that we've been able to do it so efficiently and with a growing book of business behind it.

Satya Nadella
Chairman and Chief Executive Officer

Yeah, the only thing I would add to what Amy captured was if you sort of look out, there are two things that matter, I think, and that are critical in terms of how we think about our allocation of capital, also our R&D. One is how efficient is our

planet-scale token factories? Right. I mean, that's at the end of the day what you have to do. And in order to do that, you have to start with building out a very fungible fleet. It's not like we're building one data center in one region in the world that's mega scale. We're building it out across the globe for inference, for pre-training, for post-training, for RL, for data center, or what have you. So therefore, the fungibility is super important. The second thing that we're also doing is continually modernizing the fleet. It's not like we buy one version of, say, NVIDIA and load up for all the gigawatts we have. Each year you buy, you write the Moore's Law, you continuously modernize and depreciate it. And that means you also use software to grow efficiency. I talked about, I think, 30% improvement on both serving up GPT-4.1 and 5.0, right? That's software. And by the way, it's helpful on A100s, it's helpful on GB200s, and it'll be helpful on GB300s. That's the beauty of having the efficiency of the fleet. So keep improving utilization, keep improving the efficiency. So that's what you do in the token factory. The other aspect, which Amy spoke to, is we have some of the best agent systems that matter in the high-value domains. It's in information work. That's the co-pilot system. Coding. I mean, I should also say one of the things I like about co-pilot is, I mean, co-pilot ARPUs compared to M365 ARPUs. It's expansive. The same thing that happened between server and cloud. Like we used to always say, well, is it zero sum? It turned out that the cloud was so much more expansive to the server market. The same thing is happening in AI because first you could say, hey, our ARPUs are too low. When it comes to M365, or you could say we have the opportunity with AI to be much more expansive. Same thing with tools, right? I mean, tools business was not like a leading business, whereas coding business is going to be one of the most. you know, expansive AI systems. And so we feel very good about being in that category. Same thing with security. Same thing with health. And in consumer, one of the things is it's not just about ads. It's ads plus subscriptions. That also opens up opportunity for us. So when I look at the entirety of these high-value agent systems and when we look at the efficiency and fungibility of our fleet, That's what gives us the confidence to invest both the capital and the R&D talent to go after this opportunity.

Mark Mortler
Analyst, Bernstein Research

That was pretty amazing. I really appreciate all the detail.

Jonathan Nielsen
Vice President of Investor Relations

Thanks, Mark. Operator, next question, please.

Operator
Conference Operator

The next question comes from the line of Carl Kirstead with UBS. Please proceed.

Carl Kirstead
Analyst, UBS

Okay, thank you. This one is for Amy. Amy, I certainly don't want to take you down too complex an accounting path with this question, but the investment in OpenAI that sits in other income at $4.1 billion is so large that I think the audience listening in could benefit from a little bit more color about what that is. It feels like it's so much larger than you were running through other income in prior quarters that it mustn't just be your share of the OpenAI losses. So could you just describe that and what we can expect in subsequent quarters and whether this signals any kind of accounting change. Thanks so much.

Amy Hood
Chief Financial Officer

The Q1 number was not impacted at all by the new agreement that was put in place. Let me first say that. Secondly, that increased loss was all due to our percentage of losses and opening up to the equity method. So just to be very clear. So there is not anything there that is not the increased losses from OpenAI.

Carl Kirstead
Analyst, UBS

Okay, understood. Thank you.

Jonathan Nielsen
Vice President of Investor Relations

Thanks, Carl. Thanks, Carl. Operator, next question, please.

Operator
Conference Operator

The next question comes from the line of Mark Murphy with J.P. Morgan. Please proceed.

Mark Murphy
Analyst, J.P. Morgan

Thank you so much. So we seem to be entering into a new era where the contractual commitments from a small number of AI natives are just incredibly large, not only in absolute terms, but sometimes relative to the size of the companies themselves. For instance, contracts worth hundreds of billions of dollars that are 20 times their current revenue scale. Philosophically, how do you evaluate the ability of those companies to follow through on these commitments And how do you think about placing guardrails on customer concentration for any single entity?

Satya Nadella
Chairman and Chief Executive Officer

Yeah, maybe I'll start and then Amy, you can add. I mean, it goes back a little bit, Mark, to what I said about building first the asset itself such that it's most fungible. And then to recognize the strength of even sort of our portfolio. We have a third-party business. We have a first-party business. We have third-party also spread between enterprise, digital natives. I always felt that we need a balance there because it may start with digital natives. They're always going to be the early adopters. You always have the hit app of the generation. And then essentially then it spreads throughout. The enterprise adoption cycle is just starting. And so, therefore, having over the arc of time, I think that third-party balance of customers will only increase. But it's great to have the hit first-party apps in the beginning because you can build scale that then if it's a fungible, and that's where the key is. You don't want to build for a digital native as if you were just doing hosting for them. You want to build. That's where I think some of the decision-making of ours is. It's probably getting better understood. What do we say yes to? What do we say no to? I think there was a lot of confusion. Hopefully by now anyone who switched on will figure this out. And so that's, I think, one thing we're doing on the third party. But the first party is probably where a lot of our leverage comes. And it's not even about one hit app on our first party even. Our portfolio stuff. which I just walked through in the earlier answer, gives us, again, the confidence that between that mix, we will be able to use our fleet to the maximum. And remember, these assets, especially the data centers and so on, are long assets, right? There will be many refresh cycles for any one of these when it comes to the gear. So I feel that once you think about all those dimensions, The concentration risk gets, you know, mitigated by being thoughtful about how you really ensure the build is for the broad customer base.

Amy Hood
Chief Financial Officer

And maybe just to help with another angle, Zach, I think Satya's helped a lot, is that, you know, when you think about concentration risk or delivering to any customer, you have to remember that because we're talking about this very large flexible fleet that can be used for anyone and for any purpose, 1P, 3P, and including our commercial cloud, by the way, which I should be quite clear on, it is pretty flexible in every regard. You have to remember that the CPU and GPU and

the storage gear doesn't come into play until the contracts start happening. And so you're right, some of these large contracts have delivery dates over time. So you get a lot of lead time in being able to say, oh, what's the status? So I think we're pretty thoughtful around what's always gone in our RPO balance and been considerate of that. There's always been that taken into account when we publish that booking number and publish the RPO balance.

Mark Murphy
Analyst, J.P. Morgan

Thank you very much.

Jonathan Nielsen
Vice President of Investor Relations

Thanks, Mark. Operator, next question, please.

Operator
Conference Operator

The next question comes from the line of Brad Zelnick with Deutsche Bank. Please proceed.

Brad Zelnick
Analyst, Deutsche Bank

Great. Thanks so much for taking the question, and I'll echo my congrats on an amazing start to the year. Amy, is there any way to quantify or frame the revenue impact of Azure being short on capacity? And while I appreciate the constraints you face are broad across the industry, is there risk of workloads going elsewhere, and how do you mitigate that?

Amy Hood
Chief Financial Officer

Yeah, Brad, it's a great question. You know, it's always hard to quantify precisely what would have been the revenue impact in quarter. But I would offer a way to think about it is, you know, Azure probably does bear most of the revenue impact. Because when you think about real priorities that you have to fill first, It's obviously the increasing usage and adoption and sales we've seen of M365 Copilot and the usage of Copilot Chat, which we've seen very different patterns, which we're encouraged by. It's the adoption of security features. It's the GitHub momentum. And so when you're thinking about it, that is where, and it is a priority for us, to allocate resourcing there first. And so you're all right to ask,

you know, how do I think about that We've worked very hard to try to mitigate it as best we can, but we have been short in Azure, and we've been clear on it. And I would say the other two priorities that I haven't mentioned maybe as much before is also just making sure our product teams and the AI talent that we've been able to hire into the company really over the past year and a half have access also to significant capacity because we're seeing it make the product better in a loop that is adding great benefit today into products people are using today for real-world work. And so we are making that a priority to make sure our research teams have that as well as our product engineering teams. And yes, it does impact Azure directly. That is the place where you see that prioritization. But I think it's probably hard for me to give an exact number, but it is safe to say that the number could be higher.

Keith Weiss
Analyst, Morgan Stanley

Great. Thank you.

Jonathan Nielsen
Vice President of Investor Relations

Thanks, Brad. Operator, we have time for one last question.

Operator
Conference Operator

The last question will come from the line of Kash Rangan with Goldman Sachs. Please proceed.

Kash Rangan
Analyst, Goldman Sachs

Thank you very much, Amy. I just wanted to congratulate you. I think you said before that it is possible to accelerate Azure growth while getting efficient margins, and you've done it. Congrats on that. I have one for you, Satya. With respect to the elephant in the room, following this, being a little more direct, following up on Keith Wise's question, there's talk that another hyperscaler came in and took away the business that was rightfully Microsoft's. I'm sure that there is a different point of view here. I'm wondering if you could offer some perspective on your criteria to, is it about a certain volume of business that you wish to execute on the Microsoft paper? Or is it something broader than that that I don't think maybe people fully appreciate the terminal value that Microsoft will have on its balance sheet at the end of these contracts which I think is probably being underestimated as you have a full stack and you've got the multiple vectors to monetize the databases, foundry, and to your point that you

are a platform company, not just a hyperscaler. Maybe that's what it is all about. Or maybe there's another story about you letting the other hyperscaler company come in from nowhere and claiming a big piece of that four to five-year puzzle. Thank you so much once again. Really appreciate it. And congratulations.

Satya Nadella
Chairman and Chief Executive Officer

Well, thanks, Kash. I mean, for us, again, it just always goes back to, I think, the core principle, which is build a fleet that is fungible across the planet and works for all. third party and first party and research. So that's essentially what we have done. And so when some demand comes in shapes that don't fit that goal, where it's too concentrated, not just by customer, by location, by type of skewing, right? I think Amy mentioned some very key things. When you think about the margin profile of a hyperscaler, You've got to remember there's the AI accelerator piece, but there's compute, there's storage. And so if all of the demand just comes for just one meter, that's really not a long-term business we want to be in. That's even from a third party. We have to balance it with all of our first-party stuff because that's, after all, a different margin stack for us. And then we have to fund our own R&D and model capability because in the long run, that's what's going to differentiate us. And so I look at all of those. We sort of use all of that to make sure we are saying yes to all the demand that we want. We say no to some of the demand that may be something that we could serve, but it's not in our long-term interest. And so that's sort of the decision-making we've done. And we feel very, very good about the decisions. In some sense, I feel even each time we say no to The day after, I feel better.

Amy Hood
Chief Financial Officer

And just, Cash, I think this is our last call with you, and I just want to say thanks and congratulations. It's been a privilege to work with you, and best of luck. Let me add to that. Best of luck, Cash.

Kash Rangan
Analyst, Goldman Sachs

Thanks. Thank you so much. Very kind of you.

Jonathan Nielsen
Vice President of Investor Relations

Thanks, Cash. That wraps up the Q&A portion of today's earnings call. Thank you for joining us today, and we look forward to speaking with all of you soon.

Thank you all.

Operator
Conference Operator

Thank you. This concludes today's conference. You may disconnect your lines at this time, and thank you for your participation.

Disclaimer

This conference call transcript was computer generated and almost certianly contains errors. This transcript is provided for information purposes only.**EarningsCall, LLC** makes no representation about the accuracy of the aforementioned transcript, and you are cautioned not to place undue reliance on the information provided by the transcript.